

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

A Discussion

Naganand Y

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks
 - (Shared) Filters are translation invariant

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks
 - (Shared) Filters are translation invariant
 - # Parameters independent of input size

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks
 - (Shared) Filters are translation invariant
 - # Parameters independent of input size
 - Local dependencies dominate

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks
 - (Shared) Filters are translation invariant
 - # Parameters independent of input size
 - Local dependencies dominate
- **Limitation**: Convolution, pooling defined only for regular grids

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks
 - (Shared) Filters are translation invariant
 - # Parameters independent of input size
 - Local dependencies dominate
- **Limitation**: Convolution, pooling defined only for regular grids
 - Not defined for irregular/non-Euclidean domains

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks
 - (Shared) Filters are translation invariant
 - # Parameters independent of input size
 - Local dependencies dominate
- **Limitation**: Convolution, pooling defined only for regular grids
 - Not defined for irregular/non-Euclidean domains
- **Contribution**: Define convolution, pooling on graphs

Introduction

- **CNNs**: state-of-the art on image, speech, video recognition tasks
 - (Shared) Filters are translation invariant
 - # Parameters independent of input size
 - Local dependencies dominate
- **Limitation**: Convolution, pooling defined only for regular grids
 - Not defined for irregular/non-Euclidean domains
- **Contribution**: Define convolution, pooling on graphs
 - User data on social networks
 - Gene data on biological regulatory networks
 - Log data on telecommunication networks
 - Text documents on word embeddings

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- Convolutional Neural Networks on Graphs

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs
- **Fast**

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs
- **Fast:** same complexity as CNNs

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs
- **Fast:** same complexity as CNNs
 - Linear computational complexity

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs
- **Fast:** same complexity as CNNs
 - Linear computational complexity
 - Constant learning complexity

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs
- **Fast:** same complexity as CNNs
 - Linear computational complexity
 - Constant learning complexity
- **Localized Spectral Filtering**

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs
- **Fast:** same complexity as CNNs
 - Linear computational complexity
 - Constant learning complexity
- **Localized Spectral Filtering**
 - Ideas from spectral graph theory

Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering

- **Convolutional Neural Networks on Graphs**
 - Define convolution and pooling operations on graphs
- **Fast:** same complexity as CNNs
 - Linear computational complexity
 - Constant learning complexity
- **Localized Spectral Filtering**
 - Ideas from spectral graph theory
 - Filters provable to be strictly localized in a ball of radius K

Proposed Technique

1 Design of localized convolutional filters on graphs

Proposed Technique

- 1 Design of localized convolutional filters on graphs
- 2 Graph coarsening

Proposed Technique

- 1 Design of localized convolutional filters on graphs
- 2 Graph coarsening
 - Group similar vertices together

Proposed Technique

- 1 Design of localized convolutional filters on graphs
- 2 Graph coarsening
 - Group similar vertices together
 - Build condensed, smaller graphs through these groups

Proposed Technique

- 1 Design of localized convolutional filters on graphs
- 2 Graph coarsening
 - Group similar vertices together
 - Build condensed, smaller graphs through these groups
- 3 Pooling operation

Proposed Technique

- 1 Design of localized convolutional filters on graphs
- 2 Graph coarsening
 - Group similar vertices together
 - Build condensed, smaller graphs through these groups
- 3 Pooling operation
 - Trade spatial details for higher filter details

Proposed Technique

- 1 Design of localized convolutional filters on graphs
- 2 Graph coarsening
 - Group similar vertices together
 - Build condensed, smaller graphs through these groups
- 3 Pooling operation
 - Trade spatial details for higher filter details

Notations

- **Input:** An undirected, connected, weighted $G = (V, E, W)$

Notations

- **Input**: An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$

Notations

- **Input:** An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$
- **Graph Laplacian:** $L = I_n - D^{-\frac{1}{2}} W D^{-\frac{1}{2}}$

Notations

- **Input:** An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$
- **Graph Laplacian:** $L = I_n - D^{-\frac{1}{2}} W D^{-\frac{1}{2}}$ where $D = \text{diag}(\sum_j W_{ij})$

Notations

- **Input:** An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$
- **Graph Laplacian:** $L = I_n - D^{-\frac{1}{2}} W D^{-\frac{1}{2}}$ where $D = \text{diag}(\sum_j W_{ij})$
 - Non-negative eigenvalues $\{\lambda_l\}_{l=0}^{n-1}$

Notations

- **Input:** An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$
- **Graph Laplacian:** $L = I_n - D^{-\frac{1}{2}} W D^{-\frac{1}{2}}$ where $D = \text{diag}(\sum_j W_{ij})$
 - Non-negative eigenvalues $\{\lambda_l\}_{l=0}^{n-1}$
 - Orthonormal (Fourier) basis $\{u_l\}_{l=0}^{n-1}$

Notations

- **Input:** An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$
- **Graph Laplacian:** $L = I_n - D^{-\frac{1}{2}} W D^{-\frac{1}{2}}$ where $D = \text{diag}(\sum_j W_{ij})$
 - Non-negative eigenvalues $\{\lambda_l\}_{l=0}^{n-1}$
 - Orthonormal (Fourier) basis $\{u_l\}_{l=0}^{n-1}$
 - Diagonalization $L = U \Lambda U^T$

Notations

- **Input**: An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$
- **Graph Laplacian**: $L = I_n - D^{-\frac{1}{2}} W D^{-\frac{1}{2}}$ where $D = \text{diag}(\sum_j W_{ij})$
 - Non-negative eigenvalues $\{\lambda_l\}_{l=0}^{n-1}$
 - Orthonormal (Fourier) basis $\{u_l\}_{l=0}^{n-1}$
 - Diagonalization $L = U \Lambda U^T$
- **GFT**: $\hat{x} = U^T x$

Notations

- **Input**: An undirected, connected, weighted $G = (V, E, W)$
 - **Signal** $x : V \rightarrow \mathbb{R}$ regarded as $x \in \mathbb{R}^n$
- **Graph Laplacian**: $L = I_n - D^{-\frac{1}{2}} W D^{-\frac{1}{2}}$ where $D = \text{diag}(\sum_j W_{ij})$
 - Non-negative eigenvalues $\{\lambda_l\}_{l=0}^{n-1}$
 - Orthonormal (Fourier) basis $\{u_l\}_{l=0}^{n-1}$
 - Diagonalization $L = U \Lambda U^T$
- **GFT**: $\hat{x} = U^T x$
 - Inverse, $x = U \hat{x}$

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain
- **Filtering:** A signal x is filtered by g_θ as

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain
- **Filtering:** A signal x is filtered by g_θ as

$$y = g_\theta(L)x$$

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain
- **Filtering:** A signal x is filtered by g_θ as

$$y = g_\theta(L)x = g_\theta(U\Lambda U^T)x$$

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain
- **Filtering:** A signal x is filtered by g_θ as

$$y = g_\theta(L)x = g_\theta(U\Lambda U^T)x = Ug_\theta(\Lambda)U^T x$$

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain
- **Filtering:** A signal x is filtered by g_θ as

$$y = g_\theta(L)x = g_\theta(U\Lambda U^T)x = Ug_\theta(\Lambda)U^T x$$

- Could use a non-parametric filter i.e. $g_\theta(\Lambda) = \text{diag}(\theta)$. But

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain
- **Filtering:** A signal x is filtered by g_θ as

$$y = g_\theta(L)x = g_\theta(U\Lambda U^T)x = Ug_\theta(\Lambda)U^T x$$

- Could use a non-parametric filter i.e. $g_\theta(\Lambda) = \text{diag}(\theta)$. But
 - Not localized in space

Spectral Filtering of Graph Signals

- **Convolution operator:** $a * b = U \left((U^T a) \odot (U^T b) \right)$
 - Defined on fourier domain rather than vertex domain
- **Filtering:** A signal x is filtered by g_θ as

$$y = g_\theta(L)x = g_\theta(U\Lambda U^T)x = Ug_\theta(\Lambda)U^T x$$

- Could use a non-parametric filter i.e. $g_\theta(\Lambda) = \text{diag}(\theta)$. But
 - Not localized in space
 - Learning complexitiy is $O(n)$

Polynomial Parametrization for Localized Filters

- Polynomial Filter:

$$g_{\theta}(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k$$

¹Hammond et. al, Wavelets on Graphs via Spectral Graph Theory, Applied and Computational Harmonic Analysis 2011

Polynomial Parametrization for Localized Filters

- Polynomial Filter:

$$g_{\theta}(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k$$

$$\theta \in \mathbb{R}^K$$

¹Hammond et. al, Wavelets on Graphs via Spectral Graph Theory, Applied and Computational Harmonic Analysis 2011

Polynomial Parametrization for Localized Filters

- Polynomial Filter:

$$g_{\theta}(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k$$

$$\theta \in \mathbb{R}^K$$

- ¹If minimum # edges between i and j is $> K$ then $(L^K)_{i,j} = 0$

¹Hammond et. al, Wavelets on Graphs via Spectral Graph Theory, Applied and Computational Harmonic Analysis 2011

Polynomial Parametrization for Localized Filters

- Polynomial Filter:

$$g_{\theta}(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k$$

$$\theta \in \mathbb{R}^K$$

- ¹If minimum # edges between i and j is $> K$ then $(L^K)_{i,j} = 0$
 - Filters are K -localized

¹Hammond et. al, Wavelets on Graphs via Spectral Graph Theory, Applied and Computational Harmonic Analysis 2011

Polynomial Parametrization for Localized Filters

- Polynomial Filter:

$$g_{\theta}(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k$$

$$\theta \in \mathbb{R}^K$$

- ¹If minimum # edges between i and j is $> K$ then $(L^K)_{i,j} = 0$
 - Filters are K -localized
- Learning complexity is $O(K)$ where K is support size of filter

¹Hammond et. al, Wavelets on Graphs via Spectral Graph Theory, Applied and Computational Harmonic Analysis 2011

Polynomial Parametrization for Localized Filters

- Polynomial Filter:

$$g_{\theta}(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k$$

$$\theta \in \mathbb{R}^K$$

- ¹If minimum # edges between i and j is $> K$ then $(L^K)_{i,j} = 0$
 - Filters are K -localized
- Learning complexity is $O(K)$ where K is support size of filter
 - Same as classical CNNs

¹Hammond et. al, Wavelets on Graphs via Spectral Graph Theory, Applied and Computational Harmonic Analysis 2011

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_{\theta}(\Lambda)U^T x$ costs $O(n^2)$ operations

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_{\theta}(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_{\theta}(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L
 - **Idea:** Compute polynomial recursively from L

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_{\theta}(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L
 - **Idea:** Compute polynomial recursively from L
 - Use Chebyshev expansion

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_{\theta}(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L
 - **Idea:** Compute polynomial recursively from L
 - Use Chebyshev expansion
 - $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$ with $T_0 = 1$ and $T_1 = x$

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_\theta(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L
 - **Idea:** Compute polynomial recursively from L
 - Use Chebyshev expansion
 - $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$ with $T_0 = 1$ and $T_1 = x$

- $$g_\theta(\Lambda) = \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\Lambda})$$

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_\theta(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L
 - **Idea:** Compute polynomial recursively from L
 - Use Chebyshev expansion
 - $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$ with $T_0 = 1$ and $T_1 = x$

$$g_\theta(\Lambda) = \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\Lambda})$$

- $\tilde{\Lambda} = 2\Lambda/\lambda_{\max} - I_n$

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_\theta(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L
 - **Idea:** Compute polynomial recursively from L
 - Use Chebyshev expansion
 - $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$ with $T_0 = 1$ and $T_1 = x$

$$g_\theta(\Lambda) = \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\Lambda})$$

- $\tilde{\Lambda} = 2\Lambda/\lambda_{\max} - I_n$
- Entire filtering takes $O(K|E|)$ operations

Recursive Formulation for Fast Filtering

- Multiplication with U in $y = Ug_\theta(\Lambda)U^T x$ costs $O(n^2)$ operations
- **Solution:** Exploit sparsity of L
 - **Idea:** Compute polynomial recursively from L
 - Use Chebyshev expansion
 - $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$ with $T_0 = 1$ and $T_1 = x$

$$g_\theta(\Lambda) = \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\Lambda})$$

- $\tilde{\Lambda} = 2\Lambda/\lambda_{\max} - I_n$
- Entire filtering takes $O(K|E|)$ operations
- Linear in $|E|$

Learning Filters

- Use mini-batch GD (S samples in each batch)

Learning Filters

- Use mini-batch GD (S samples in each batch)
 - j -th feature of sample s is

$$y_{s,j} = \sum_{i=1}^{F_{in}} g_{\theta_{i,j}}(L) x_{s,i}$$

Learning Filters

- Use mini-batch GD (S samples in each batch)
 - j -th feature of sample s is

$$y_{s,j} = \sum_{i=1}^{F_{in}} g_{\theta_{i,j}}(L) x_{s,i}$$

- Backprop requires

$$\frac{\partial J}{\partial \theta_{i,j}} = \sum_{s=1}^S T_k(\tilde{L}) \frac{\partial J}{\partial y_{s,j}}$$

Learning Filters

- Use mini-batch GD (S samples in each batch)
 - j -th feature of sample s is

$$y_{s,j} = \sum_{i=1}^{F_{in}} g_{\theta_{i,j}}(L) x_{s,i}$$

- Backprop requires

$$\frac{\partial J}{\partial \theta_{i,j}} = \sum_{s=1}^S T_k(\tilde{L}) \frac{\partial J}{\partial y_{s,j}}$$

and

$$\frac{\partial J}{\partial x_{s,i}} = \sum_{j=1}^{F_{out}} g_{\theta_{i,j}}(L) \frac{\partial J}{\partial y_{s,j}}$$

Learning Filters

- Use mini-batch GD (S samples in each batch)
 - j -th feature of sample s is

$$y_{s,j} = \sum_{i=1}^{F_{in}} g_{\theta_{i,j}}(L) x_{s,i}$$

- Backprop requires

$$\frac{\partial J}{\partial \theta_{i,j}} = \sum_{s=1}^S T_k(\tilde{L}) \frac{\partial J}{\partial y_{s,j}}$$

and

$$\frac{\partial J}{\partial x_{s,i}} = \sum_{j=1}^{F_{out}} g_{\theta_{i,j}}(L) \frac{\partial J}{\partial y_{s,j}}$$

- Requires $O(K|E|F_{in}F_{out}S)$ operations

Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs

Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs
 - Cluster similar vertices together

Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs
 - Cluster similar vertices together
 - But graph clustering is NP-hard

Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs
 - Cluster similar vertices together
 - But graph clustering is NP-hard
- Graclus graph clustering software

Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs
 - Cluster similar vertices together
 - But graph clustering is NP-hard
- Graclus graph clustering software
 - Uses a greedy algorithm

Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs
 - Cluster similar vertices together
 - But graph clustering is NP-hard
- Graclus graph clustering software
 - Uses a greedy algorithm
 - Pick an unmarked i and match with one of its unmarked neighbours j that maximizes normalized cut

$$W_{ij} \left(\frac{1}{d_i} + \frac{1}{d_j} \right)$$

Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs
 - Cluster similar vertices together
 - But graph clustering is NP-hard
- Graclus graph clustering software
 - Uses a greedy algorithm
 - Pick an unmarked i and match with one of its unmarked neighbours j that maximizes normalized cut

$$W_{ij} \left(\frac{1}{d_i} + \frac{1}{d_j} \right)$$

- Repeat until all nodes explored

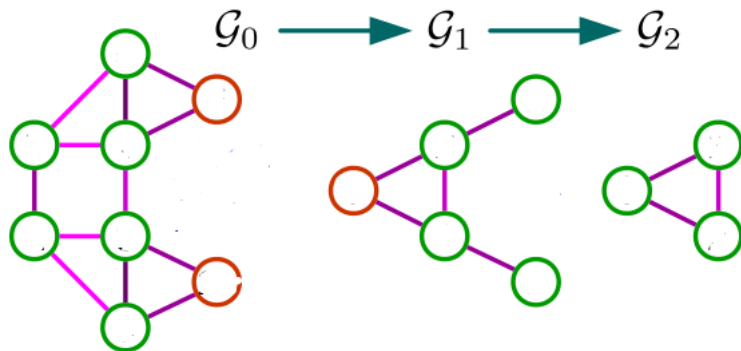
Graph Coarsening

- Pooling requires meaningful neighbourhoods on graphs
 - Cluster similar vertices together
 - But graph clustering is NP-hard
- Graclus graph clustering software
 - Uses a greedy algorithm
 - Pick an unmarked i and match with one of its unmarked neighbours j that maximizes normalized cut

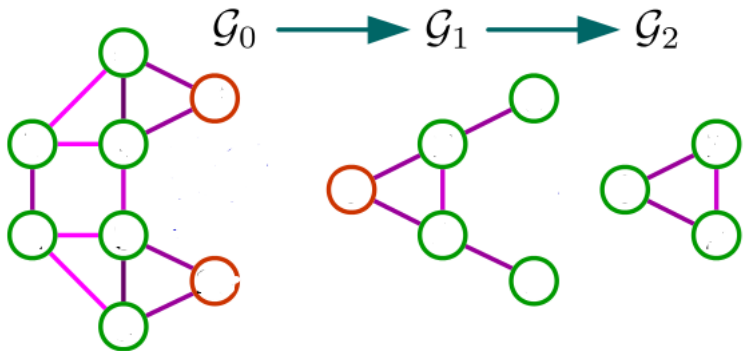
$$W_{ij} \left(\frac{1}{d_i} + \frac{1}{d_j} \right)$$

- Repeat until all nodes explored
- Divides $\#$ nodes by approximately 2 (there may exist a few singletons, non-matched nodes)

Graph Coarsening: Example

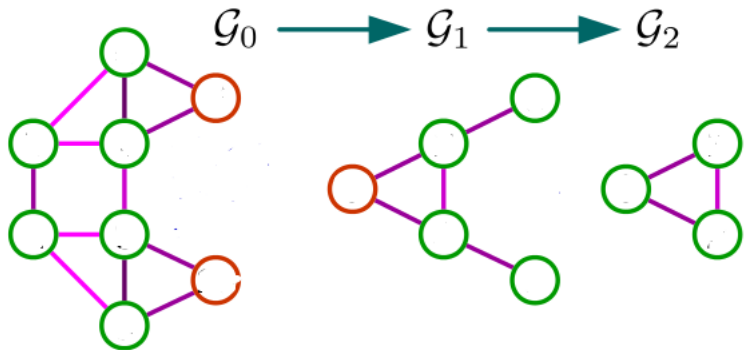


Graph Coarsening: Example



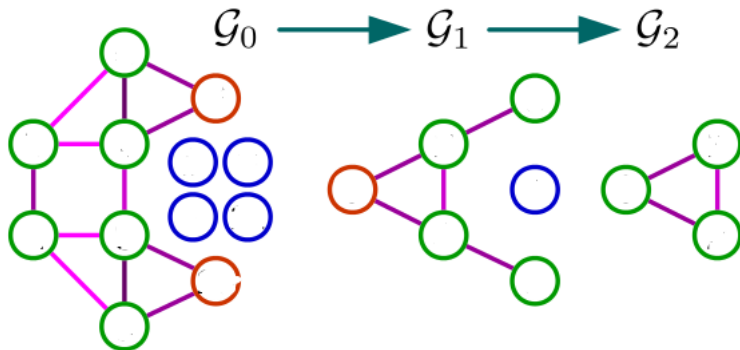
Graculus produces \mathcal{G}_1 and \mathcal{G}_2 on \mathcal{G}_0 as input

Max Pooling: Example



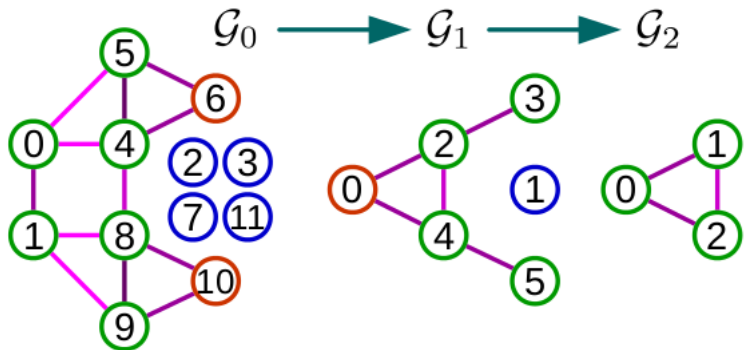
Suppose max pooling of size 4 needs to be carried out (Note: $x \in \mathbb{R}^8$)

Max Pooling: Example



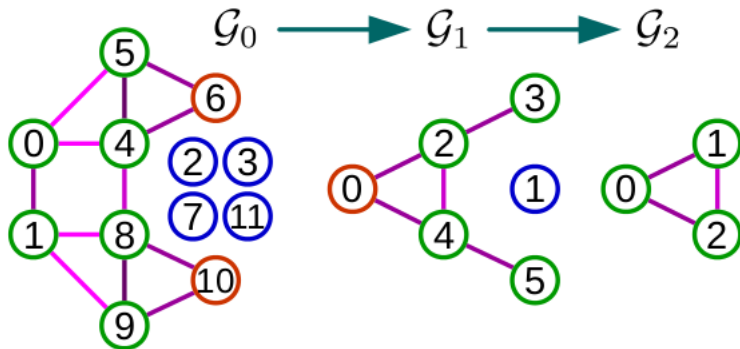
Add fake nodes to pair with singleton nodes

Max Pooling: Example



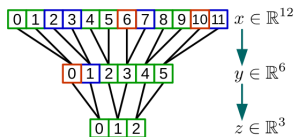
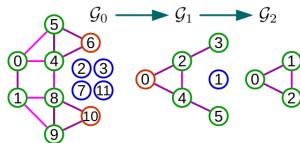
Nodes in V_2 : ordered arbitrarily, nodes in V_1 , V_0 : ordered consequently

Max Pooling: Example

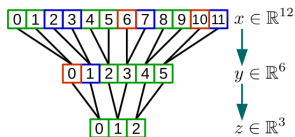
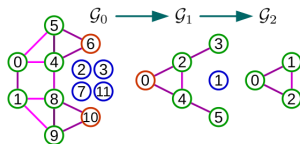


Observation: Node k has $2k$ and $2k + 1$ as children

Max Pooling: Example

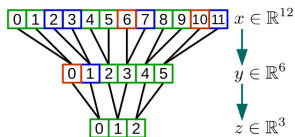
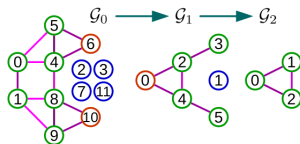


Max Pooling: Example



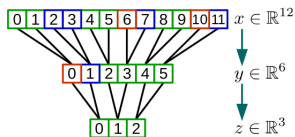
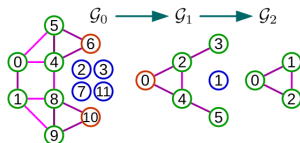
- $z = \left(\max(x_0, x_1), \max(x_4, x_5, x_6), \max(x_8, x_9, x_{10}) \right) \in \mathbb{R}^3$

Max Pooling: Example



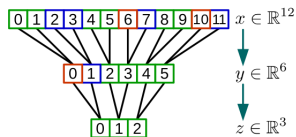
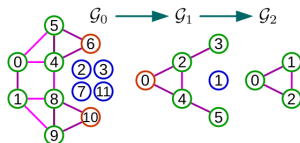
- $z = \left(\max(x_0, x_1), \max(x_4, x_5, x_6), \max(x_8, x_9, x_{10}) \right) \in \mathbb{R}^3$
 - x_2, x_3, x_7, x_{11} are set to a neutral value

Max Pooling: Example



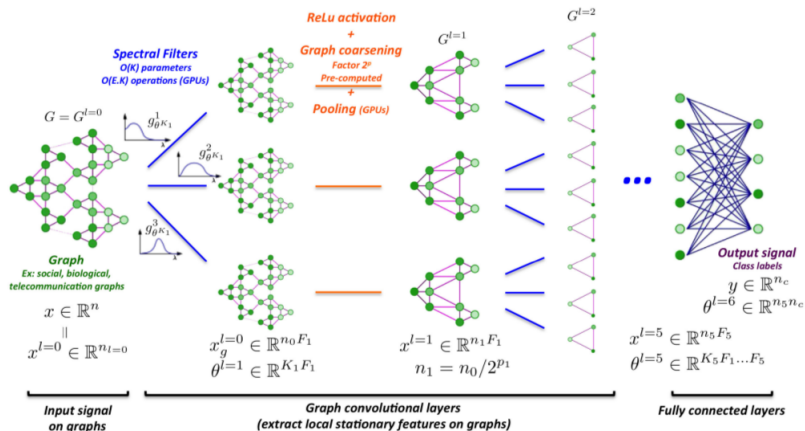
- $z = \left(\max(x_0, x_1), \max(x_4, x_5, x_6), \max(x_8, x_9, x_{10}) \right) \in \mathbb{R}^3$
 - x_2, x_3, x_7, x_{11} are set to a neutral value
 - Analogous to pooling a regular 1D signal

Max Pooling: Example



- $z = \left(\max(x_0, x_1), \max(x_4, x_5, x_6), \max(x_8, x_9, x_{10}) \right) \in \mathbb{R}^3$
 - x_2, x_3, x_7, x_{11} are set to a neutral value
 - Analogous to pooling a regular 1D signal
 - Very efficient, satisfies parallel architectures such as GPUs as memory accesses are local

Architecture



Experiments: MNIST as a Sanity Check

- Construct an 8-NN graph

Experiments: MNIST as a Sanity Check

- Construct an 8-NN graph
 - 70000 digits on 2D grid becomes graph with 976 ($28^2 + 192$) nodes

Experiments: MNIST as a Sanity Check

- Construct an 8-NN graph
 - 70000 digits on 2D grid becomes graph with 976 ($28^2 + 192$) nodes
 - Weights of k -NN similarity graph are

$$W_{ij} = \exp\left(-\frac{\|z_i - z_j\|_2^2}{\sigma^2}\right)$$

z_i is 2D coordinate of pixel i

Experiments: MNIST as a Sanity Check

- Construct an 8-NN graph
 - 70000 digits on 2D grid becomes graph with 976 ($28^2 + 192$) nodes
 - Weights of k -NN similarity graph are

$$W_{ij} = \exp\left(-\frac{\|z_i - z_j\|_2^2}{\sigma^2}\right)$$

z_i is 2D coordinate of pixel i

Model	Architecture	Accuracy
Classical CNN	C32-P4-C64-P4-FC512	99.33
Proposed graph CNN	GC32-P4-GC64-P4-FC512	99.14

- Ck : Convolutional layer k feature maps
- GCk : Graph convolutional layer k feature maps
- FCk : Fully connected layer k hidden units
- Pk : Pooling layer size, stride k

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract 10k most common words from 93k+ unique words

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract $10k$ most common words from $93k+$ unique words
 - Represent doc x by bag-of-words model (normalized across words)

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract $10k$ most common words from $93k+$ unique words
 - Represent doc x by bag-of-words model (normalized across words)
 - Construct a 16-NN graph with $W_{ij} = \exp(-\frac{\|z_i - z_j\|_2^2}{\sigma^2})$ where z_i is word2vec embedding of word i

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract $10k$ most common words from $93k+$ unique words
 - Represent doc x by bag-of-words model (normalized across words)
 - Construct a 16-NN graph with $W_{ij} = \exp(-\frac{\|z_i - z_j\|_2^2}{\sigma^2})$ where z_i is word2vec embedding of word i

Model	Accuracy
Linear SVM	65.90
Multinomial Naive Bayes	68.51
Softmax	66.28
FC2500	64.64
FC2500-FC500	65.76
GC32	68.26

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract 10k most common words from 93k+ unique words
 - Represent doc x by bag-of-words model (normalized across words)
 - Construct a 16-NN graph with $W_{ij} = \exp(-\frac{\|z_i - z_j\|_2^2}{\sigma^2})$ where z_i is word2vec embedding of word i

Model	Accuracy
Linear SVM	65.90
Multinomial Naive Bayes	68.51
Softmax	66.28
FC2500	64.64
FC2500-FC500	65.76
GC32	68.26

- All models trained 20 epochs by Adam optimizer (initial learning rate = 0.001)

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract 10k most common words from 93k+ unique words
 - Represent doc x by bag-of-words model (normalized across words)
 - Construct a 16-NN graph with $W_{ij} = \exp(-\frac{\|z_i - z_j\|_2^2}{\sigma^2})$ where z_i is word2vec embedding of word i

Model	Accuracy
Linear SVM	65.90
Multinomial Naive Bayes	68.51
Softmax	66.28
FC2500	64.64
FC2500-FC500	65.76
GC32	68.26

- All models trained 20 epochs by Adam optimizer (initial learning rate = 0.001)
- Support of GC32 is $K = 5$

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract 10k most common words from 93k+ unique words
 - Represent doc x by bag-of-words model (normalized across words)
 - Construct a 16-NN graph with $W_{ij} = \exp(-\frac{\|z_i - z_j\|_2^2}{\sigma^2})$ where z_i is word2vec embedding of word i

Model	Accuracy
Linear SVM	65.90
Multinomial Naive Bayes	68.51
Softmax	66.28
FC2500	64.64
FC2500-FC500	65.76
GC32	68.26

- All models trained 20 epochs by Adam optimizer (initial learning rate = 0.001)
- Support of GC32 is $K = 5$
- Proposed model does not outperform multinomial Bayes but defeats FC networks which require much more parameters

Text Categorization on 20NEWS

- To demonstrate versatility of model to work with unstructured data
 - Extract 10k most common words from 93k+ unique words
 - Represent doc x by bag-of-words model (normalized across words)
 - Construct a 16-NN graph with $W_{ij} = \exp(-\frac{\|z_i - z_j\|_2^2}{\sigma^2})$ where z_i is word2vec embedding of word i

Model	Accuracy
Linear SVM	65.90
Multinomial Naive Bayes	68.51
Softmax	66.28
FC2500	64.64
FC2500-FC500	65.76
GC32	68.26

- All models trained 20 epochs by Adam optimizer (initial learning rate = 0.001)
- Support of GC32 is $K = 5$
- Proposed model does not outperform multinomial Bayes but defeats FC networks which require much more parameters

Comparison of Spectral Filters

- The first formulation² of graph CNN defined filter as $g_{\theta}(\Lambda) = B\theta$

²Bruna et. al, Spectral Networks and Deep Locally Connected Networks on Graphs [17/19](#)

Comparison of Spectral Filters

- The first formulation² of graph CNN defined filter as $g_{\theta}(\Lambda) = B\theta$

²Bruna et. al, Spectral Networks and Deep Locally Connected Networks on Graphs [17/19](#)

Comparison of Spectral Filters

- The first formulation² of graph CNN defined filter as $g_\theta(\Lambda) = B\theta$ where $B \in \mathbb{R}^{n \times K}$ is cubic B-spline basis

²Bruna et. al, Spectral Networks and Deep Locally Connected Networks on Graphs [17/19](#)

Comparison of Spectral Filters

- The first formulation² of graph CNN defined filter as $g_{\theta}(\Lambda) = B\theta$ where $B \in \mathbb{R}^{n \times K}$ is cubic B-spline basis

Dataset	Architecture	Accuracy		
		Non-Param	Spline	Chebyshev
MNIST	GC10	95.75	97.26	97.48
MNIST	GC32-P4-GC64-P4-FC512	96.28	97.15	99.14

²Bruna et. al, Spectral Networks and Deep Locally Connected Networks on Graphs [17/19](#)

Comparison of Spectral Filters

- The first formulation² of graph CNN defined filter as $g_{\theta}(\Lambda) = B\theta$ where $B \in \mathbb{R}^{n \times K}$ is cubic B-spline basis

Dataset	Architecture	Accuracy		
		Non-Param	Spline	Chebyshev
MNIST	GC10	95.75	97.26	97.48
MNIST	GC32-P4-GC64-P4-FC512	96.28	97.15	99.14

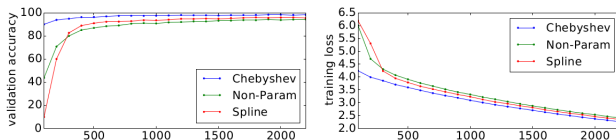


Figure 4: Plots of validation accuracy and training loss for the first 2000 iterations on MNIST.

²Bruna et. al, Spectral Networks and Deep Locally Connected Networks on Graphs [17/19](#)

Comparison of Computational Efficiency

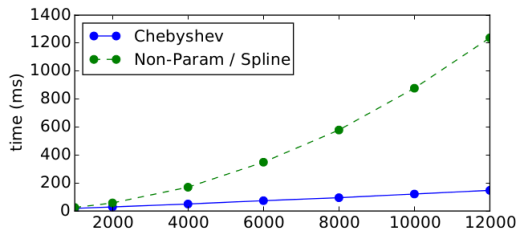


Figure 3: Time to process a mini-batch of $S = 100$ 20NEWS documents w.r.t. the number of words n .

Comparison of Computational Efficiency

• $O(n)$ vs. $O(n^2)$

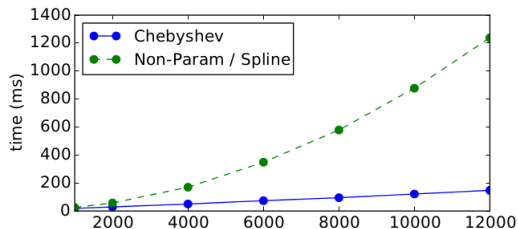


Figure 3: Time to process a mini-batch of $S = 100$ 20NEWS documents w.r.t. the number of words n .

Comparison of Computational Efficiency

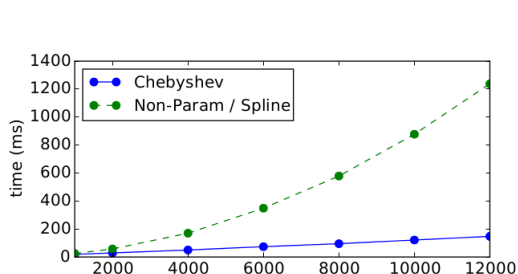
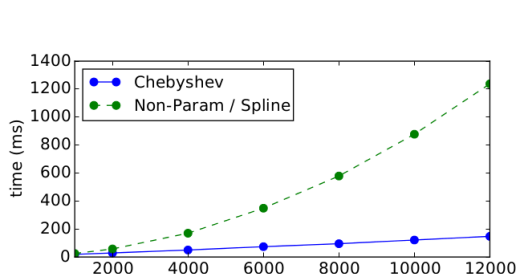


Figure 3: Time to process a mini-batch of $S = 100$ 20NEWS documents w.r.t. the number of words n .

- $O(n)$ vs. $O(n^2)$
- Measured runtime is total training time divided by #gradient steps

Comparison of Computational Efficiency



- $O(n)$ vs. $O(n^2)$
- Measured runtime is total training time divided by #gradient steps

Figure 3: Time to process a mini-batch of $S = 100$ 20NEWS documents w.r.t. the number of words n .

Model	Architecture	Time (ms)		
		CPU	GPU	Speedup
Classical CNN	C32-P4-C64-P4-FC512	210	31	6.77x
Proposed graph CNN	GC32-P4-GC64-P4-FC512	1600	200	8.00x

Table 4: Time to process a mini-batch of $S = 100$ MNIST images.

Conclusion

- Introduced mathematical and computational foundations of efficient generalization of CNNs on graphs

Conclusion

- Introduced mathematical and computational foundations of efficient generalization of CNNs on graphs
- Introduced model whose computational complexity linear with dimensionality of data

Conclusion

- Introduced mathematical and computational foundations of efficient generalization of CNNs on graphs
- Introduced model whose computational complexity linear with dimensionality of data
- Experiments have shown ability of model to extract local, stationary features through convolutional layers