

Learning Robust Representations of Text

A Discussion

Naganand Y

Summary

- Deep nets are sensitive to noise, adversarial attacks
- Present regularization method to limit network sensitivity to inputs
 - Models become more robust
 - Ideas are inspired by computer vision
- Achieve superior performance on noisy inputs, out-of-domain data on sentiment datasets

Introduction

- Primary cause of neural nets' vulnerability is linear nature¹
 - LSTMs, ReLUs, maxout designed linearly to facilitate optimization
- Fawzi et al.² showed linear models not robust to adversarial noise
- Present a regularization method to make neural nets more robust to noise
 - Inspired by Rifai et al ³

¹Goodfellow et al. , Explaining and harnessing adversarial examples, ICLR 2014

²Analysis of classifiers' robustness to adversarial perturbations

³Contractive autoencoders: Explicit invariance during feature extraction, ICML 2011 3/7

Approach

- **Intuition:** Minimize ability of features to perturb predictions
 - to stabilize predictions
- **Idea:** Train models using first-order derivatives of training loss as part of regularization term
 - Necessitates second-order derivatives for computing gradient

Training for Robustness

- **Training:** SGD to min L (measures $y_{pred} - y_{true}$)
 - w : Input, a sequence of (discrete) words
 - h : fixed-size vector of continuous values representing w
 - $y_{pred} = f(h)$
- **Goal:** Learn models that are more robust to strange/invalid inputs
 - y_{pred} remains stable on perturbations on w (or h)
- **Application:** Transfer learning scenarios such as domain adaptation
 - Inputs in distinct domains drawn from different distributions
 - Highly variable but convey same information
 - Different word choice, different syntactic structures, typographical errors, stylistic changes, etc

Robust Regularization

- Minimize variation of output when noise applied to input

- $\Delta_y = f(x + \Delta_x) - f(x)$

-

$$\lim_{\Delta_x \rightarrow 0} \Delta_y = \lim_{\Delta_x \rightarrow 0} (f(x + \Delta_x) - f(x)) = \frac{\partial y}{\partial x} \cdot \Delta_x$$

- Minimising noise sensitivity \equiv minimising $\left\| \frac{\partial y}{\partial x} \right\|_F$

- $\mathcal{L} = L + \lambda \cdot \left\| \frac{\partial L}{\partial h} \right\|_2$

- Supports gradient optimization
 - Need to compute second-order derivatives of L during back-propagation

Experiments and Datasets

- Model f is CNN proposed by Yoon Kim⁴
 - MR: Sentence polarity dataset
 - Subj: Subjectivity dataset
 - CR: Customer review dataset
 - SST: Stanford Sentiment Treebank, using the 3-class configuration
- **Noise**: Apply world-level dropout noise to each document
 - Randomly replace words by a unique sentinel symbol
 - Apply this to each word with probability $\alpha \in \{0, 0.1, 0.2, 0.3\}$
- Cross-domain evaluation
 - Train on one dataset, apply it to another
 - Pair MR (movie reviews) and CR (product reviews) that use same label set

¹Convolutional neural networks for sentence classification, EMNLP 2014